

```
In [33]: !pip install numpy pandas matplotlib seaborn plotly requests tqdm opencv-python
```

Looking in indexes: <https://pypi.tuna.tsinghua.edu.cn/simple>  
Requirement already satisfied: numpy in /environment/miniconda3/lib/python3.10/site-packages (1.24.1)  
Requirement already satisfied: pandas in /environment/miniconda3/lib/python3.10/site-packages (2.1.2)  
Requirement already satisfied: matplotlib in /environment/miniconda3/lib/python3.10/site-packages (3.8.1)  
Requirement already satisfied: seaborn in /environment/miniconda3/lib/python3.10/site-packages (0.13.0)  
Requirement already satisfied: plotly in /environment/miniconda3/lib/python3.10/site-packages (5.19.0)  
Requirement already satisfied: requests in /environment/miniconda3/lib/python3.10/site-packages (2.31.0)  
Requirement already satisfied: tqdm in /environment/miniconda3/lib/python3.10/site-packages (4.65.0)  
Requirement already satisfied: opencv-python in /environment/miniconda3/lib/python3.10/site-packages (4.8.1.78)  
Requirement already satisfied: pillow in /environment/miniconda3/lib/python3.10/site-packages (9.3.0)  
Requirement already satisfied: wandb in /environment/miniconda3/lib/python3.10/site-packages (0.16.3)  
Requirement already satisfied: python-dateutil<=2.8.2 in /environment/miniconda3/lib/python3.10/site-packages (from pandas) (2.8.2)  
Requirement already satisfied: pytz<=2020.1 in /environment/miniconda3/lib/python3.10/site-packages (from pandas) (2023.3.post1)  
Requirement already satisfied: tzdata<=2022.1 in /environment/miniconda3/lib/python3.10/site-packages (from pandas) (2023.3)  
Requirement already satisfied: contourpy<=1.0.1 in /environment/miniconda3/lib/python3.10/site-packages (from matplotlib) (1.2.0)  
Requirement already satisfied: cycler<=0.10 in /environment/miniconda3/lib/python3.10/site-packages (from matplotlib) (0.12.1)  
Requirement already satisfied: fonttools<=4.22.0 in /environment/miniconda3/lib/python3.10/site-packages (from matplotlib) (4.44.0)  
Requirement already satisfied: kiwisolver<=1.3.1 in /environment/miniconda3/lib/python3.10/site-packages (from matplotlib) (1.4.5)  
Requirement already satisfied: packaging<=20.0 in /environment/miniconda3/lib/python3.10/site-packages (from matplotlib) (23.0)  
Requirement already satisfied: pyparsing<=2.3.1 in /environment/miniconda3/lib/python3.10/site-packages (from matplotlib) (3.1.1)  
Requirement already satisfied: tenacity<=6.2.0 in /environment/miniconda3/lib/python3.10/site-packages (from plotly) (8.2.3)  
Requirement already satisfied: charset-normalizer<4,>=2 in /environment/miniconda3/lib/python3.10/site-packages (from requests) (2.0.4)  
Requirement already satisfied: idna<4,>=2.5 in /environment/miniconda3/lib/python3.10/site-packages (from requests) (2.10)  
Requirement already satisfied: urllib3<3,>=1.21.1 in /environment/miniconda3/lib/python3.10/site-packages (from requests) (2.2.1)  
Requirement already satisfied: certifi<=2017.4.17 in /environment/miniconda3/lib/python3.10/site-packages (from requests) (2023.7.22)  
Requirement already satisfied: Click!=8.0.0,>=7.1 in /environment/miniconda3/lib/python3.10/site-packages (from wandb) (7.1.2)  
Requirement already satisfied: GitPython!=3.1.29,>=1.0.0 in /environment/miniconda3/lib/python3.10/site-packages (from wandb) (3.1.42)  
Requirement already satisfied: psutil<=5.0.0 in /environment/miniconda3/lib/python3.10/site-packages (from wandb) (5.9.5)  
Requirement already satisfied: sentry-sdk<=1.0.0 in /environment/miniconda3/lib/python3.10/site-packages (from wandb) (1.40.5)  
Requirement already satisfied: docker-pycreds<=0.4.0 in /environment/miniconda3/lib/python3.10/site-packages (from wandb) (0.4.0)  
Requirement already satisfied: PyYAML in /environment/miniconda3/lib/python3.10/s

```
ite-packages (from wandb) (6.0.1)
Requirement already satisfied: setproctitle in /environment/miniconda3/lib/python
3.10/site-packages (from wandb) (1.3.3)
Requirement already satisfied: setuptools in /environment/miniconda3/lib/python3.
10/site-packages (from wandb) (67.8.0)
Requirement already satisfied: appdirs>=1.4.3 in /environment/miniconda3/lib/pyth
on3.10/site-packages (from wandb) (1.4.4)
Requirement already satisfied: protobuf!=4.21.0,<5,>=3.19.0 in /environment/minic
onda3/lib/python3.10/site-packages (from wandb) (4.23.4)
Requirement already satisfied: six>=1.4.0 in /environment/miniconda3/lib/python3.
10/site-packages (from docker-pycreds>=0.4.0->wandb) (1.16.0)
Requirement already satisfied: gitdb<5,>=4.0.1 in /environment/miniconda3/lib/pyt
hon3.10/site-packages (from GitPython!=3.1.29,>=1.0.0->wandb) (4.0.11)
Requirement already satisfied: smmap<6,>=3.0.1 in /environment/miniconda3/lib/pyt
hon3.10/site-packages (from gitdb<5,>=4.0.1->GitPython!=3.1.29,>=1.0.0->wandb)
(5.0.1)
```

## Download and install Pytorch

```
In [34]: !pip3 install torch torchvision torchaudio --extra-index-url https://download.py
```

Looking in indexes: <https://pypi.tuna.tsinghua.edu.cn/simple>, <https://download.pytorch.org/whl/cu113>  
Requirement already satisfied: torch in /environment/miniconda3/lib/python3.10/site-packages (2.0.1+cu118)  
Requirement already satisfied: torchvision in /environment/miniconda3/lib/python3.10/site-packages (0.15.2+cu118)  
Requirement already satisfied: torchaudio in /environment/miniconda3/lib/python3.10/site-packages (2.0.2+cu118)  
Requirement already satisfied: filelock in /environment/miniconda3/lib/python3.10/site-packages (from torch) (3.9.0)  
Requirement already satisfied: typing-extensions in /environment/miniconda3/lib/python3.10/site-packages (from torch) (4.8.0)  
Requirement already satisfied: sympy in /environment/miniconda3/lib/python3.10/site-packages (from torch) (1.11.1)  
Requirement already satisfied: networkx in /environment/miniconda3/lib/python3.10/site-packages (from torch) (3.0)  
Requirement already satisfied: Jinja2 in /environment/miniconda3/lib/python3.10/site-packages (from torch) (3.1.2)  
Requirement already satisfied: triton==2.0.0 in /environment/miniconda3/lib/python3.10/site-packages (from torch) (2.0.0)  
Requirement already satisfied: cmake in /environment/miniconda3/lib/python3.10/site-packages (from triton==2.0.0->torch) (3.25.0)  
Requirement already satisfied: lit in /environment/miniconda3/lib/python3.10/site-packages (from triton==2.0.0->torch) (15.0.7)  
Requirement already satisfied: numpy in /environment/miniconda3/lib/python3.10/site-packages (from torchvision) (1.24.1)  
Requirement already satisfied: requests in /environment/miniconda3/lib/python3.10/site-packages (from torchvision) (2.31.0)  
Requirement already satisfied: pillow!=8.3.\*,>=5.3.0 in /environment/miniconda3/lib/python3.10/site-packages (from torchvision) (9.3.0)  
Requirement already satisfied: MarkupSafe>=2.0 in /environment/miniconda3/lib/python3.10/site-packages (from Jinja2->torch) (2.1.2)  
Requirement already satisfied: charset-normalizer<4,>=2 in /environment/miniconda3/lib/python3.10/site-packages (from requests->torchvision) (2.0.4)  
Requirement already satisfied: idna<4,>=2.5 in /environment/miniconda3/lib/python3.10/site-packages (from requests->torchvision) (2.10)  
Requirement already satisfied: urllib3<3,>=1.21.1 in /environment/miniconda3/lib/python3.10/site-packages (from requests->torchvision) (2.2.1)  
Requirement already satisfied: certifi>=2017.4.17 in /environment/miniconda3/lib/python3.10/site-packages (from requests->torchvision) (2023.7.22)  
Requirement already satisfied: mpmath>=0.19 in /environment/miniconda3/lib/python3.10/site-packages (from sympy->torch) (1.2.1)

In [35]: !wget https://zihao-openmmlab.obs.cn-east-3.myhuaweicloud.com/20220716-mmclassif

```
--2024-02-26 15:04:33-- https://zihao-openmmlab.obs.cn-east-3.myhuaweicloud.com/20220716-mmclassification/dataset/SimHei.ttf
Connecting to 172.16.0.13:5848... connected.
Proxy request sent, awaiting response... 200 OK
Length: 10050868 (9.6M) [application/x-font-ttf]
Saving to: 'SimHei.ttf.1'
```

```
SimHei.ttf.1          100%[=====] 9.58M  20.9MB/s   in 0.5s
```

```
2024-02-26 15:04:34 (20.9 MB/s) - 'SimHei.ttf.1' saved [10050868/10050868]
```

## Create a catalogue

```
In [36]: import os
```

```
In [37]: # Store the results file
# os.mkdir('output')

# Store the trained model weights
os.mkdir('checkpoint')

# Store the generated charts
os.mkdir('diagrams')
```

## Setting matplotlib Chinese and English fonts

```
In [38]: ## Font Environment Settings
import matplotlib.pyplot as plt
from matplotlib import rcParams
from matplotlib.font_manager import FontProperties

# global font settings
SimSun = FontProperties(fname='/home/featurize/SimHei.ttf') # Used to display C
plt.rcParams['axes.unicode_minus'] = False # Used to display the negative sign
Times_New_Roman = FontProperties(fname='/home/featurize/times.ttf')

# mixed font settings
config = {
    "font.family": 'serif',
    "font.size": 80,
    "mathtext.fontset": 'stix',
    "font.serif": ['SimSun'],
}
rcParams.update(config)

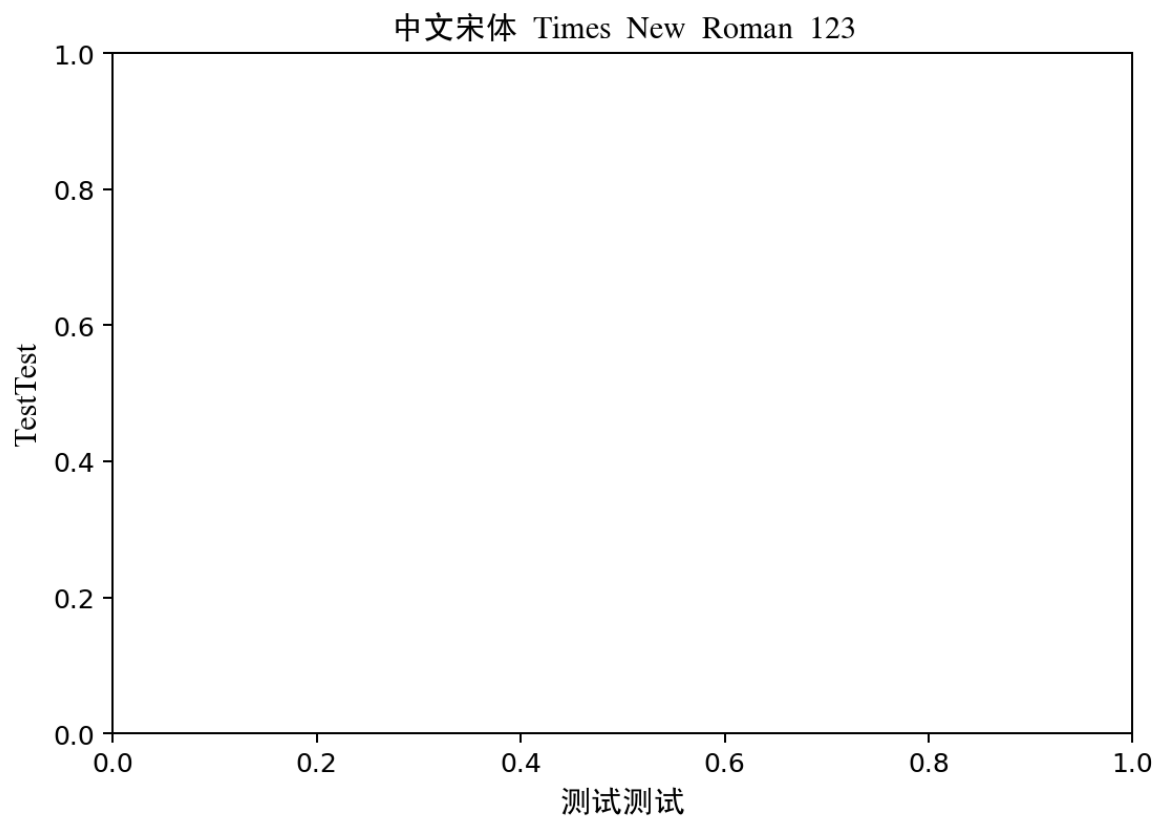
#Canvas Settings
fig = plt.figure(num=1, figsize=(9, 6), dpi=180)
ax = plt.axes((0.23,0.23,0.6,0.6))

# Application of font effects
ax.set_title('中文宋体  $\mathrm{Times}$   $\mathrm{New}$   $\mathrm{Roman}$   $\mathrm{}$ , fontproperties=SimSun, fontsize=12)

ax.set_xlabel('测试测试', fontproperties=SimSun, fontsize=12)

ax.set_ylabel('TestTest', fontproperties=Times_New_Roman, fontsize=12)

plt.show()
```



```
In [42]: !sudo snap install tree
```

```
snap "tree" is already installed, see 'snap help refresh'
```

```
In [17]: !tree /home/featurize/data -L 2
```

```
/home/featurize/data
├── fruit25_split.zip
├── train
│   ├── CherryTomatoes
│   ├── Mangosteen
│   ├── MomordicaCharantia
│   ├── NavelOrange
│   ├── Sandsugaroranges
│   ├── apple
│   ├── banana
│   ├── carrot
│   ├── cherries
│   ├── cucumber
│   ├── durian
│   ├── grape
│   ├── hamimelon
│   ├── kiwi
│   ├── lemon
│   ├── lichee
│   ├── longan
│   ├── mango
│   ├── pear
│   ├── pineapple
│   ├── pitaya
│   ├── pomegranate
│   ├── strawberry
│   ├── tomato
│   └── watermelon
└── val
    ├── Cherrytomatoes
    ├── Mangosteen
    ├── MomordicaCharantia
    ├── NavelOrange
    ├── Sandsugaroranges
    ├── apple
    ├── banana
    ├── carrot
    ├── cherries
    ├── cucumber
    ├── durian
    ├── grape
    ├── hamimelon
    ├── kiwi
    ├── lemon
    ├── lichee
    ├── longan
    ├── mango
    ├── pear
    ├── pineapple
    ├── pitaya
    ├── pomegranate
    ├── strawberry
    ├── tomato
    └── watermelon
```

52 directories, 1 file

```
In [19]: import time
import os
```

```

import numpy as np
from tqdm import tqdm

import torch
import torchvision
import torch.nn as nn
import torch.nn.functional as F

import matplotlib.pyplot as plt
%matplotlib inline

import warnings
warnings.filterwarnings("ignore")

```

```

In [20]: # test cpu
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print('device', device)

```

device cuda:0

```

In [21]: from torchvision import transforms

# Training Set Image Preprocessing - RCTN: Scaling, Cropping, Turn Tensor, Normalise
train_transform = transforms.Compose([transforms.RandomResizedCrop(224),
                                     transforms.RandomHorizontalFlip(),
                                     transforms.ToTensor(),
                                     transforms.Normalize([0.485, 0.456, 0.406])
                                     ])

# Test Set Image Preprocessing - RCTN: Scaling, Cropping, Turn Tensor, Normalise
test_transform = transforms.Compose([transforms.Resize(256),
                                    transforms.CenterCrop(224),
                                    transforms.ToTensor(),
                                    transforms.Normalize(
                                        mean=[0.485, 0.456, 0.406],
                                        std=[0.229, 0.224, 0.225])
                                    ])

```

```

In [37]: # Dataset folder path
dataset_dir = '/home/featurize/data'

```

```

In [ ]:

```

```

In [70]: train_path = os.path.join(dataset_dir, 'train')
test_path = os.path.join(dataset_dir, 'val')
print('Training_set_path', train_path)
print('Testing_set_path', test_path)

```

Training\_set\_path /home/featurize/data/train  
Testing\_set\_path /home/featurize/data/val

```

In [79]: from torchvision import datasets

# Load training set
train_dataset = datasets.ImageFolder(train_path, train_transform)

# Load Test Set
test_dataset = datasets.ImageFolder(test_path, test_transform)

```



```
In [71]: print('Number of images in the training set', len(train_dataset))
print('Number of categories', len(train_dataset.classes))
print('Name of each category', train_dataset.classes)
```

Number of images in the training set 3649

Number of categories 25

Name of each category ['CherryTomatoes', 'Mangosteen', 'MomordicaCharantia', 'NavelOrange', 'Sandsugaroranges', 'apple', 'banana', 'carrot', 'cherries', 'cucumber', 'durian', 'grape', 'hamimelon', 'kiwi', 'lemon', 'lichee', 'longan', 'mango', 'pear', 'pineapple', 'pitaya', 'pomegranate', 'strawberry', 'tomato', 'watermelon']

```
In [72]: print('Number of test set images', len(test_dataset))
print('Number of categories', len(test_dataset.classes))
print('Name of each category', test_dataset.classes)
```

Number of test set images 898

Number of categories 25

Name of each category ['Cherrytomatoes', 'Mangosteen', 'MomordicaCharantia', 'NavelOrange', 'Sandsugaroranges', 'apple', 'banana', 'carrot', 'cherries', 'cucumber', 'durian', 'grape', 'hamimelon', 'kiwi', 'lemon', 'lichee', 'longan', 'mango', 'pear', 'pineapple', 'pitaya', 'pomegranate', 'strawberry', 'tomato', 'watermelon']

```
In [73]: # Name of each category
class_names = train_dataset.classes
n_class = len(class_names)
```

```
In [74]: class_names
```

```
Out[74]: ['CherryTomatoes',
'Mangosteen',
'MomordicaCharantia',
'NavelOrange',
'Sandsugaroranges',
'apple',
'banana',
'carrot',
'cherries',
'cucumber',
'durian',
'grape',
'hamimelon',
'kiwi',
'lemon',
'lichee',
'longan',
'mango',
'pear',
'pineapple',
'pitaya',
'pomegranate',
'strawberry',
'tomato',
'watermelon']
```

```
In [75]: # Mapping relationship: category to index number
train_dataset.class_to_idx
```

```
Out[75]: {'CherryTomatoes': 0,
          'Mangosteen': 1,
          'MomordicaCharantia': 2,
          'NavelOrange': 3,
          'Sandsugaroranges': 4,
          'apple': 5,
          'banana': 6,
          'carrot': 7,
          'cherries': 8,
          'cucumber': 9,
          'durian': 10,
          'grape': 11,
          'hamimelon': 12,
          'kiwi': 13,
          'lemon': 14,
          'lichee': 15,
          'longan': 16,
          'mango': 17,
          'pear': 18,
          'pineapple': 19,
          'pitaya': 20,
          'pomegranate': 21,
          'strawberry': 22,
          'tomato': 23,
          'watermelon': 24}
```

```
In [76]: # Mapping relationship: index number to category
idx_to_labels = {y:x for x,y in train_dataset.class_to_idx.items()}
```

```
In [77]: idx_to_labels
```

```
Out[77]: {0: 'CherryTomatoes',
          1: 'Mangosteen',
          2: 'MomordicaCharantia',
          3: 'NavelOrange',
          4: 'Sandsugaroranges',
          5: 'apple',
          6: 'banana',
          7: 'carrot',
          8: 'cherries',
          9: 'cucumber',
          10: 'durian',
          11: 'grape',
          12: 'hamimelon',
          13: 'kiwi',
          14: 'lemon',
          15: 'lichee',
          16: 'longan',
          17: 'mango',
          18: 'pear',
          19: 'pineapple',
          20: 'pitaya',
          21: 'pomegranate',
          22: 'strawberry',
          23: 'tomato',
          24: 'watermelon'}
```

```
In [78]: # Save as local npy file
np.save('idx_to_labels.npy', idx_to_labels)
```

```
np.save('labels_to_idx.npy', train_dataset.class_to_idx)
```

## Define the data loader DataLoader

```
In [52]: from torch.utils.data import DataLoader
```

```
In [53]: BATCH_SIZE = 32

# Data Loader for the training set
train_loader = DataLoader(train_dataset,
                          batch_size=BATCH_SIZE,
                          shuffle=True,
                          num_workers=4
                        )

# Data Loader for Test Sets
test_loader = DataLoader(test_dataset,
                        batch_size=BATCH_SIZE,
                        shuffle=False,
                        num_workers=4
                      )
```

## View images and annotations for a batch

```
In [54]: # DataLoader is a python generator that returns a batch of data per call.
images, labels = next(iter(train_loader))
```

```
In [55]: images.shape
```

```
Out[55]: torch.Size([32, 3, 224, 224])
```

```
In [56]: labels
```

```
Out[56]: tensor([16, 16, 10,  1, 24,  3,  4,  4, 17, 23,  3,  3, 15,  8, 19, 12,  1, 10,
                  19,  4, 18,  3, 18,  4,  5, 16, 24,  0,  8, 20, 24, 11])
```


## Visualising a batch of images and annotations

```
In [57]: # Converting the Tensor tensor in a dataset to numpy's array data type
images = images.numpy()
```

```
In [58]: images[5].shape
```

```
Out[58]: (3, 224, 224)
```


```
In [59]: plt.hist(images[5].flatten(), bins=50)
plt.show()
```



```
In [60]: # Preprocessed images in batch
idx = 2
plt.imshow(images[idx].transpose((1,2,0))) # 转为(224, 224, 3)
plt.title('label:'+str(labels[idx].item()))
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

```
Out[60]: Text(0.5, 1.0, 'label:10')
```



```
In [61]: label = labels[idx].item()
```

```
In [62]: label
```


```
Out[62]: 10
```

```
In [63]: pred_classname = idx_to_labels[label]
```

```
In [64]: pred_classname
```

```
Out[64]: 'durian'
```

```
In [65]: # original image
idx = 2
mean = np.array([0.485, 0.456, 0.406])
std = np.array([0.229, 0.224, 0.225])
plt.imshow(np.clip(images[idx].transpose((1,2,0)) * std + mean, 0, 1))
plt.title('label:' + pred_classname)
plt.show()
```



## Toolkit to be used for importing training

```
In [66]: from torchvision import models
import torch.optim as optim
```

Randomly initialise all weights of the model  
and train all layers from scratch

```
In [68]: model = models.resnet18(pretrained=False) # Load only the model structure, not t

model.fc = nn.Linear(model.fc.in_features, n_class)

optimizer = optim.Adam(model.parameters())
```

## Training configuration

```
In [111... model = model.to(device)

# Cross Entropy Loss Function
criterion = nn.CrossEntropyLoss()

# Training rounds Epoch
EPOCHS = 20
```

```
In [81]: # Get a batch of data and annotations
images, labels = next(iter(train_loader))
images = images.to(device)
labels = labels.to(device)
```

```
In [82]: # Input model to perform forward prediction
outputs = model(images)
```

```
In [83]: # Get the predicted category logit scores for all images in the current batch
outputs.shape
```

```
Out[83]: torch.Size([32, 25])
```

```
In [84]: # From Logit, calculate the average cross-entropy loss function for each sample
loss = criterion(outputs, labels)
```

```
In [85]: optimizer.zero_grad() # Clearing the gradient
loss.backward() # backward propagation
optimizer.step() # Optimisation Updates
```

```
In [86]: # Get the prediction categories for all images in the current batch
_, preds = torch.max(outputs, 1)
```

```
In [87]: preds
```

```
Out[87]: tensor([13, 14, 19,  3, 14,  9,  3, 19, 10, 14, 19, 19, 14, 13, 12, 21,  3, 19,
                15, 14, 14, 19,  3, 14,  3, 15, 19, 13,  3, 13, 14,  3],
                device='cuda:0')
```

```
In [88]: labels
```

```
Out[88]: tensor([23, 15, 10, 24, 15,  3,  9, 15, 17,  0,  5,  1, 24,  0, 21, 14,  6, 13,
                24, 19,  0, 10, 10,  0, 19, 10, 23,  2, 12, 22,  1,  2],
                device='cuda:0')
```

## Run the full training

```
In [94]: # Iterate through each EPOCH
for epoch in tqdm(range(EPOCHS)):

    model.train()

    for images, labels in train_loader: # Get a batch of the training set with
        images = images.to(device)
        labels = labels.to(device)

        outputs = model(images) # Forward Prediction, get the prediction
        loss = criterion(outputs, labels) # Compare the predictions with the annotations

        optimizer.zero_grad()
        loss.backward() # Loss function back propagation of neural network
        optimizer.step() # Optimisation to update neural network
```

100%|██████████| 20/20 [01:46<00:00, 5.33s/it]

## Initial testing on a test set

```
In [95]: model.eval()
with torch.no_grad():
    correct = 0
    total = 0
    for images, labels in tqdm(test_loader): # Get a batch of the test set with
        images = images.to(device)
        labels = labels.to(device)
        outputs = model(images) # Forward prediction, get the prediction
        _, preds = torch.max(outputs, 1) # Obtain the category corresponding to the prediction
        total += labels.size(0)
        correct += (preds == labels).sum() # Number of correct samples predicted

    print('The accuracy on the test set is {:.3f} %'.format(100 * correct / total))
```

100%|██████████| 29/29 [00:01<00:00, 17.24it/s]

The accuracy on the test set is 62.695 %

```
In [114... torch.save(model, 'checkpoint/fruit25_pytorch_xintian.pth')
```